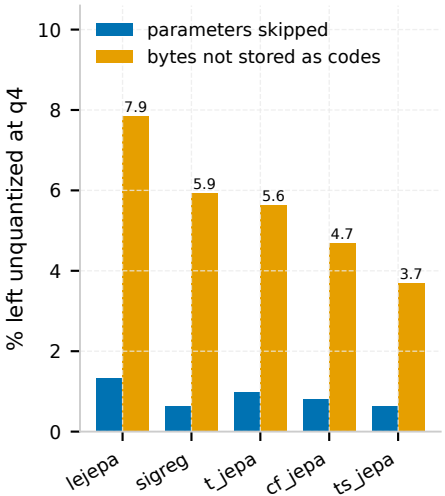
Coverage of affine quantization
parameters skipped
bytes not stored as codes
% left unquantized at q4
0
2
4
6
8
10
7.9
5.9
5.6
4.7
3.7
lejepa
sigreg
t_jepa
cf_jepa
ts_jepa

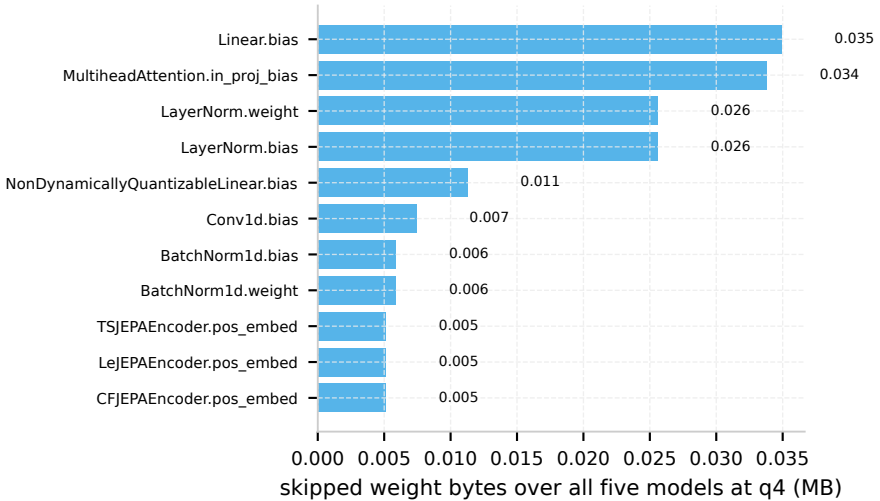
Skipped tensor classes
Linear.bias
MultiheadAttention.in_proj_bias
LayerNorm.weight
LayerNorm.bias
NonDynamicallyQuantizableLinear.bias
Conv1d.bias
BatchNorm1d.bias
BatchNorm1d.weight
TSJEPAEncoder.pos_embed
LeJEPAEncoder.pos_embed
CFJEPAEncoder.pos_embed
0.035
0.034
0.026
0.026
0.011
0.007
0.006
0.006
0.005
0.005
0.005
0.000
0.005
0.010
0.015
0.020
0.025
0.030
0.035
skipped weight bytes over all five models at q4 (MB)